

# **Linear Algebra**

## **Problem Set 2 - Solutions**



**Abolfazl Ranjbar**

Course Instructors: Dr. Elham Ghasroddashti and Dr. Peyman Adibi

Computer Engineering Department  
University of Isfahan

Fall Semester 2024

-۱-  $S$  را یک زیر فضای از  $V$  می نامند هرگاه،

- $\forall s, t \in S \rightarrow s + t \in S$
  - $\forall s \in S, \forall a \in F \rightarrow as \in S$

نشان دهید، در فضای برداری دو بعدی  $\mathbb{R}^2$  هر خط راستی که از مبدأ عبور کند، یک زیر فضای  
برداری از  $\mathbb{R}^2$  است.

$$S = \{(x, y) \in \mathbb{R}^2 : ax + by = 0\}$$

بستہ بودن نسبت بہ جمیں ہے

$$\begin{cases} s = (x_1, y_1) \in S \\ t = (x_2, y_2) \in S \end{cases} \rightarrow s + t \in S$$

$\rightsquigarrow (x_1 + x_2, y_1 + y_2) \in S$

$$\text{we know that } \begin{cases} ax_1 + bx_1 = 0 \\ ax_2 + bx_2 = 0 \end{cases} \rightarrow a(x_1 + x_2) + b(y_1 + y_2) = 0 + 0 = 0$$

بسم الله الرحمن الرحيم

If  $s = (x_1, y_1) \in S$ ,  $c \in R \xrightarrow{\text{so}} cs \in S \hookrightarrow (cx_1, cy_1) \in S$

$$a(cx_1) + b(cy_1) = c(ax_1) + c(by_1) = c(ax_1 + by_1) = 0$$

$$\rightarrow a(cx_i) + b(cy_i) = 0 \rightarrow (cx, cy) \in S$$

$\sin R^2$  is always positive  $\Rightarrow$

آیا در فضای برداری دو بعدی  $\mathbb{R}^2$  هر خط راستی که از مبدأ عبور نکند، یک زیر فضای برداری از  $\mathbb{R}^2$  است؟

$$S = \{(x, y) \in \mathbb{R}^2 : ax + by = k\}$$

$$\text{if } (x_1, y_1), (x_2, y_2) \in S \rightarrow (x_1 + x_2, y_1 + y_2) \in S$$

$$\left\{ \begin{array}{l} (x_1, y_1) \in S \rightarrow ax_1 + \frac{by_1}{x_1} = k \\ (x_2, y_2) \in S \rightarrow ax_2 + \frac{by_2}{x_2} = K \end{array} \right. \rightarrow a(x_1 + x_2) + b\left(\frac{y_1 + y_2}{x_1 + x_2}\right) = (ax_1 + \frac{by_1}{x_1}) + (ax_2 + \frac{by_2}{x_2}) \rightarrow = K + K = 2K$$

عندما راسخ في ذهنه أن صياغة مفهوم زیر خدمتی بخطاب  $R^2$  نتیجه

$$n_1 + b_1 y_1 \rightarrow = K + K = 2K$$

$\rightarrow K \neq 2K \rightarrow$

۲- استقلال خطی یا وابستگی خطی بردارهای زیر را بررسی کنید.

$$\alpha_1 u_1 + \alpha_2 u_2 + \alpha_3 u_3 = \alpha_1 \begin{bmatrix} -2 \\ 1 \end{bmatrix} + \alpha_2 \begin{bmatrix} -1 \\ -3 \end{bmatrix} + \alpha_3 \begin{bmatrix} 4 \\ -2 \end{bmatrix} \rightsquigarrow \text{هموار برابر غیر متناسب ندارد}$$

$$\alpha_1 = 2 \rightsquigarrow \begin{cases} -\alpha_2 + 4\alpha_3 = 4 \\ -3\alpha_2 - 2\alpha_3 = -2 \end{cases} \rightsquigarrow -5\alpha_2 = 0 \rightarrow \alpha_2 = 0, \alpha_3 = 0 \rightsquigarrow \text{وابسته خطی}$$

$$u_1 = \begin{bmatrix} 1 \\ -2 \end{bmatrix}, u_2 = \begin{bmatrix} -1 \\ 3 \end{bmatrix}, u_3 = \begin{bmatrix} 1 \\ -2 \end{bmatrix} \quad (\text{پ})$$

$$\begin{array}{l} \left. \begin{array}{l} \alpha_1 - \alpha_2 + \alpha_3 = 0 \\ -2\alpha_1 + 3\alpha_2 + \alpha_3 = 0 \end{array} \right\} \rightsquigarrow 3\alpha_1 - 4\alpha_2 = 0 \\ \left. \begin{array}{l} 3\alpha_1 + 4\alpha_2 - 2\alpha_3 = 0 \\ 4\alpha_1 + 2\alpha_2 - 2\alpha_3 = 0 \end{array} \right\} \rightsquigarrow 7\alpha_1 + 2\alpha_2 = 0 \end{array} \rightsquigarrow 18\alpha_1 = 0$$

$$\alpha_1 = 0 \rightsquigarrow \alpha_2 = 0 \rightsquigarrow \alpha_3 = 0 \rightsquigarrow \text{مستقل خلی}$$

- ۳- آیا مجموعه زیر برای فضای برداری مورد نظر تشکیل یک پایه می دهد

$$(\mathfrak{N}^3) \text{ فضای برداری } v_1 = \begin{bmatrix} 1 \\ -1 \\ 1 \end{bmatrix}, v_2 = \begin{bmatrix} -1 \\ 2 \\ -2 \end{bmatrix}, v_3 = \begin{bmatrix} -1 \\ 4 \\ -4 \end{bmatrix} \quad (\text{الف})$$

$$\alpha_1 v_1 + \alpha_2 v_2 + \alpha_3 v_3 = 0 \left\{ \begin{array}{l} \alpha_1 - \alpha_2 - \alpha_3 = 0 \\ -\alpha_1 + 2\alpha_2 + 4\alpha_3 = 0 \\ \alpha_1 - 2\alpha_2 - 4\alpha_3 = 0 \end{array} \right\} \stackrel{\Leftrightarrow}{=} \alpha_2 + 3\alpha_3 = 0 \left\{ \begin{array}{l} \alpha_2 = 3 \\ \alpha_3 = -1 \end{array} \right.$$

وابسته خطی سه گزینه  $\alpha_1 = 2$

۴- بردارهای زیر را در نظر بگیرید. با استفاده از الگوریتم گرام-اشمیت بردارهای پایه یکه متعامد را بدست آورید.

$$\tilde{f}_k = \vec{a}_k - \sum_{i=1}^{k-1} (f_i^T \vec{a}_k) \vec{f}_i$$

$$u_1 = \begin{bmatrix} 1 \\ 2 \\ 1 \\ 0 \end{bmatrix}, u_2 = \begin{bmatrix} 3 \\ 3 \\ 0 \end{bmatrix}, u_3 = \begin{bmatrix} 2 \\ 0 \\ 0 \end{bmatrix}, u_4 = \begin{bmatrix} -10 \\ 1 \\ -6 \end{bmatrix}$$

$$f_1 = \frac{\tilde{f}_1}{\|\tilde{f}_1\|} = \frac{u_1}{\|u_1\|} = \begin{bmatrix} \frac{1}{\sqrt{6}} \\ \frac{2}{\sqrt{6}} \\ \frac{1}{\sqrt{6}} \\ 0 \end{bmatrix}$$

$$\tilde{f}_2 = u_2 - (f_1^T u_2) \vec{f}_1 = \begin{bmatrix} 3 \\ 3 \\ 0 \end{bmatrix} - 2\sqrt{6} \begin{bmatrix} \frac{1}{\sqrt{6}} \\ \frac{2}{\sqrt{6}} \\ \frac{1}{\sqrt{6}} \end{bmatrix} = \begin{bmatrix} 3 \\ 3 \\ 0 \end{bmatrix} - \begin{bmatrix} 2 \\ 4 \\ 0 \end{bmatrix} = \begin{bmatrix} 1 \\ -1 \\ 0 \end{bmatrix} \quad f_2 = \begin{bmatrix} \frac{1}{\sqrt{3}} \\ -\frac{1}{\sqrt{3}} \\ \frac{1}{\sqrt{3}} \\ 0 \end{bmatrix}$$

$$\begin{aligned} \tilde{f}_3 &= u_3 - (f_2^T u_3) \vec{f}_2 - (f_1^T u_3) \vec{f}_1 = \begin{bmatrix} 2 \\ -10 \\ 0 \end{bmatrix} - (4\sqrt{3}) \begin{bmatrix} \frac{1}{\sqrt{3}} \\ -\frac{1}{\sqrt{3}} \\ \frac{1}{\sqrt{3}} \end{bmatrix} + (3\sqrt{6}) \begin{bmatrix} \frac{1}{\sqrt{6}} \\ \frac{2}{\sqrt{6}} \\ \frac{1}{\sqrt{6}} \end{bmatrix} \\ &= \begin{bmatrix} \frac{2}{\sqrt{10}} \\ -\frac{4}{\sqrt{10}} \\ 0 \end{bmatrix} - \begin{bmatrix} 4 \\ 0 \\ 0 \end{bmatrix} + \begin{bmatrix} 3 \\ 0 \\ 0 \end{bmatrix} = \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix} \quad \sim f_3 = \frac{1}{\sqrt{2}} \begin{bmatrix} 1 \\ 0 \\ -1 \end{bmatrix} = \begin{bmatrix} \frac{1}{\sqrt{2}} \\ 0 \\ -\frac{1}{\sqrt{2}} \end{bmatrix} \end{aligned}$$

$$\begin{aligned} \tilde{f}_4 &= u_4 - (f_3^T u_4) \vec{f}_3 - (f_2^T u_4) \vec{f}_2 - (f_1^T u_4) \vec{f}_1 \\ &= \begin{bmatrix} -2 \\ 1 \\ 2 \end{bmatrix} - (2\sqrt{2}) \begin{bmatrix} \frac{1}{\sqrt{2}} \\ 0 \\ 0 \end{bmatrix} + (3\sqrt{2}) \begin{bmatrix} \frac{1}{\sqrt{3}} \\ -\frac{1}{\sqrt{3}} \\ \frac{1}{\sqrt{3}} \end{bmatrix} + (\sqrt{6}) \begin{bmatrix} \frac{1}{\sqrt{6}} \\ \frac{2}{\sqrt{6}} \\ \frac{1}{\sqrt{6}} \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix} \quad \sim f_4 = \begin{bmatrix} 0 \\ 0 \\ 1 \end{bmatrix} \end{aligned}$$

۵- اگر فرایند گرام اشمیت را به یک دسته بردار وابسته خطی اعمال نماییم. نتیجه نهایی چه خواهد شد؟ با انتخاب سه بردار وابسته خطی پاسخ خود را بررسی نمایید.

$$\text{example: } u_1 = \begin{bmatrix} 0 \\ 0 \\ 1 \end{bmatrix}, u_2 = \begin{bmatrix} 0 \\ 1 \\ 0 \end{bmatrix}, u_3 = \begin{bmatrix} 0 \\ 1 \\ 2 \end{bmatrix}$$

$$f_1 = \frac{\vec{u}_1}{\|u_1\|} = \begin{bmatrix} 0 \\ 0 \\ 1 \end{bmatrix} \quad \text{کوثر بطریه صفر نداشته.}$$

$$\tilde{f}_2 = u_2 - (f_1^T u_2) \vec{f}_1 = \begin{bmatrix} 0 \\ 1 \\ 0 \end{bmatrix} - \begin{bmatrix} 0 \\ 1 \\ 0 \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix} \rightarrow f_2 = \frac{\vec{f}_2}{\|\vec{f}_2\|} = \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix} \quad \sim$$

- 4.1** Minimizing mean square distance to a set of vectors. Let  $x_1, \dots, x_L$  be a collection of  $n$ -vectors. In this exercise you will fill in the missing parts of the argument to show that the vector  $z$  which minimizes the sum-square distance to the vectors,

$$J(z) = \|x_1 - z\|^2 + \dots + \|x_L - z\|^2,$$

is the average or centroid of the vectors,  $\bar{x} = (1/L)(x_1 + \dots + x_L)$ . (This result is used in one of the steps in the  $k$ -means algorithm. But here we have simplified the notation.)

- (a) Explain why, for any  $z$ , we have

$$J(z) = \sum_{i=1}^L \|x_i - \bar{x} - (z - \bar{x})\|^2 \stackrel{\text{apply}}{=} \sum_{i=1}^L (\|x_i - \bar{x}\|^2 - 2(x_i - \bar{x})^T(z - \bar{x})) + L\|z - \bar{x}\|^2.$$

The norm squared expanded

- (b) Explain why  $\sum_{i=1}^L (x_i - \bar{x})^T(z - \bar{x}) = 0$ . Hint. Write the left-hand side as

$$\left( \sum_{i=1}^L (x_i - \bar{x}) \right)^T (z - \bar{x}), \underbrace{\left( \sum_{i=1}^L x_i - L\bar{x} \right)}_{\cancel{L\bar{x}}} \underbrace{(z - \bar{x})}_{\cancel{so \rightarrow = 0}} = 0$$

and argue that the left-hand vector is 0.

- (c) Combine the results of (a) and (b) to get  $J(z) = \sum_{i=1}^L \|x_i - \bar{x}\|^2 + L\|z - \bar{x}\|^2$ . Explain why for any  $z \neq \bar{x}$ , we have  $J(z) > J(\bar{x})$ . This shows that the choice  $z = \bar{x}$  minimizes  $J(z)$ .

$$J(z) = \sum_{i=1}^L \|x_i - \bar{x}\|^2 + L\|z - \bar{x}\|^2 > \sum_{i=1}^L \|x_i - \bar{x}\|^2 = J(\bar{x})$$

$$\begin{aligned} J(z) &= \sum_{i=1}^L (\|x_i - \bar{x}\|^2 - 2(x_i - \bar{x})^T(z - \bar{x}) + L\|z - \bar{x}\|^2) \\ &= \sum_{i=1}^L (\|x_i - \bar{x}\|^2) - 2 \sum_{i=1}^L ((x_i - \bar{x})^T(z - \bar{x})) + L\|z - \bar{x}\|^2 \end{aligned}$$

$$\sum_{i=1}^L (x_i - \bar{x})^T(z - \bar{x}) = 0$$

$$J(z) = \sum_{i=1}^L (\|x_i - \bar{x}\|^2) + L\|z - \bar{x}\|^2$$



**4.2** *k-means with nonnegative, proportions, or Boolean vectors.* Suppose that the vectors  $x_1, \dots, x_N$  are clustered using  $k$ -means, with group representatives  $z_1, \dots, z_k$ .

- (a) Suppose the original vectors  $x_i$  are nonnegative, i.e., their entries are nonnegative. Explain why the representatives  $z_j$  are also nonnegative.
- (b) Suppose the original vectors  $x_i$  represent proportions, i.e., their entries are nonnegative and sum to one. (This is the case when  $x_i$  are word count histograms, for example.) Explain why the representatives  $z_j$  also represent proportions, i.e., their entries are nonnegative and sum to one.
- (c) Suppose the original vectors  $x_i$  are Boolean, i.e., their entries are either 0 or 1. Give an interpretation of  $(z_j)_i$ , the  $i$ th entry of the  $j$  group representative.

*Hint.* Each representative is the average of some of the original vectors.

(a) 
$$z_j = \frac{1}{|G_j|} \sum_{k \in G_j} x_k \quad \text{if } x \text{ are nonnegative}$$

$\hookrightarrow z_j \text{ is nonnegative too}$

(b) 
$$z_j^T z_j = \frac{1}{|G_j|} \sum_{k \in G_j} z_j^T x_k = \frac{|G_j|}{|G_j|} = 1$$

(c) because  $z_j$  is a fraction of the vectors in the group  $\rightarrow$  if  $z_j = 1$   
 then all the vectors in that group are 1 or if it close to 1  $\rightarrow$  most of  
 vectors are one.

**4.3** *Linear separation in 2-way partitioning.* Clustering a collection of vectors into  $k = 2$  groups is called 2-way partitioning, since we are partitioning the vectors into 2 groups, with index sets  $G_1$  and  $G_2$ . Suppose we run  $k$ -means, with  $k = 2$ , on the  $n$ -vectors  $x_1, \dots, x_N$ . Show that there is a nonzero vector  $w$  and a scalar  $v$  that satisfy

$$w^T x_i + v \geq 0 \text{ for } i \in G_1, \quad w^T x_i + v \leq 0 \text{ for } i \in G_2.$$

In other words, the affine function  $f(x) = w^T x + v$  is greater than or equal to zero on the first group, and less than or equal to zero on the second group. This is called **linear separation** of the two groups (although **affine separation** would be more accurate).

*Hint.* Consider the function  $\|x - z_1\|^2 - \|x - z_2\|^2$ , where  $z_1$  and  $z_2$  are the group representatives.

$$\|x_i - z_2\|^2 - \|x_i - z_1\|^2 \geq 0$$

$$\|x_i - z_2\|^2 - \|x_i - z_1\|^2 \leq 0$$

expanding : 
$$\|x_i\|^2 - 2x_i^T z_2 + \|z_2\|^2 - \|x_i\|^2 + 2x_i^T z_1 - \|z_1\|^2$$
  

$$= -2x_i^T (z_2 - z_1) + \|z_2\|^2 - \|z_1\|^2$$

format 
$$= -2(z_2 - z_1)^T x_i + \|z_2\|^2 - \|z_1\|^2$$

$\downarrow M \quad \downarrow \sigma$

**4.4 Pre-assigned vectors.** Suppose that some of the vectors  $x_1, \dots, x_N$  are assigned to specific groups. For example, we might insist that  $x_{27}$  be assigned to group 5. Suggest a simple modification of the  $k$ -means algorithm that respects this requirement. Describe a practical example where this might arise, when each vector represents  $n$  features of a medical patient.

→ first we override the partitioning step, then we enforce the given assignment instead of allowing the algorithm to resign it.  
 The remaining vectors are assigned to groups as usual based on their position

**5.1 Linear independence of stacked vectors.** Consider the stacked vectors

$$c_1 = \begin{bmatrix} a_1 \\ b_1 \end{bmatrix}, \dots, c_k = \begin{bmatrix} a_k \\ b_k \end{bmatrix},$$

where  $a_1, \dots, a_k$  are  $n$ -vectors and  $b_1, \dots, b_k$  are  $m$ -vectors.

- (a) Suppose  $a_1, \dots, a_k$  are linearly independent. (We make no assumptions about the vectors  $b_1, \dots, b_k$ .) Can we conclude that the stacked vectors  $c_1, \dots, c_k$  are linearly independent?
- (b) Now suppose that  $a_1, \dots, a_k$  are linearly dependent. (Again, with no assumptions about  $b_1, \dots, b_k$ .) Can we conclude that the stacked vectors  $c_1, \dots, c_k$  are linearly dependent?

stack vector:  
 refers to vectors formed  
 with concatenating two or  
 more vectors { vertically = column  
 horizontally = row vector }

(a) yes, suppose:  $\alpha^T C = 0 \rightarrow \alpha_1 c_1 + \dots + \alpha_n c_n = 0$   
 $\rightarrow \alpha_1 \begin{bmatrix} a_1 \\ b_1 \end{bmatrix} + \dots + \alpha_n \begin{bmatrix} a_n \\ b_n \end{bmatrix} = 0$

as you can see  $\rightarrow \alpha_1 a_1 + \dots + \alpha_n a_n = 0 \rightarrow$  is linear independent

$\rightarrow \alpha_1, \dots, \alpha_n = 0 \rightarrow$  then adding "b" vector doesn't change anything.

(b) nope!, let's use the example above:

$\alpha_1 a_1 + \dots + \alpha_n a_n = 0 \rightarrow$  is linear independent

$\rightarrow$  so there is some  $\alpha_i$  that are not zero

$\hookrightarrow$  so by adding "b" vector, those coefficients might not work

$$\rightarrow \alpha_1 \beta_1 + \dots + \alpha_n \beta_n \neq 0$$

- 5.2** A surprising discovery. An intern at a quantitative hedge fund examines the daily returns of around 400 stocks over one year (which has 250 trading days). She tells her supervisor that she has discovered that the returns of one of the stocks, Google (GOOG), can be expressed as a linear combination of the others, which include many stocks that are unrelated to Google (say, in a different type of business or sector).

Her supervisor then says: "It is overwhelmingly unlikely that a linear combination of the returns of unrelated companies can reproduce the daily return of GOOG. So you've made a mistake in your calculations."

Is the supervisor right? Did the intern make a mistake? Give a very brief explanation.

→ The intern is correct, with 400 stocks in a 250-dimensional space (trading days), linear dependence guarantees that at least one stock's returns can be expressed as linear combination of the others.

- 5.4** Norm of linear combination of orthonormal vectors. Suppose  $a_1, \dots, a_k$  are orthonormal  $n$ -vectors, and  $x = \beta_1 a_1 + \dots + \beta_k a_k$ , where  $\beta_1, \dots, \beta_k$  are scalars. Express  $\|x\|$  in terms of  $\beta = (\beta_1, \dots, \beta_k)$ .

$$\|x\|^2 = x^T x = (\beta_1 \alpha_1 + \dots + \beta_k \alpha_k)^T (\beta_1 \alpha_1 + \dots + \beta_k \alpha_k) =$$

as  $\alpha_i$  is orthonormal  $\rightarrow \alpha_i^T \alpha_i = 1$

$$\text{so } \rightarrow \|x\|^2 = \beta_1^2 + \dots + \beta_k^2 = \|\beta\|^2 \rightarrow \|x\| = \|\beta\|$$

- 5.5** Orthogonalizing vectors. Suppose that  $a$  and  $b$  are any  $n$ -vectors. Show that we can always find a scalar  $\gamma$  so that  $(a - \gamma b) \perp b$ , and that  $\gamma$  is unique if  $b \neq 0$ . (Give a formula for the scalar  $\gamma$ .) In other words, we can always subtract a multiple of a vector from another one, so that the result is orthogonal to the original vector. The orthogonalization step in the Gram-Schmidt algorithm is an application of this.

$$\begin{aligned} \rightarrow (a - \gamma b)^T b = 0 &\rightarrow a^T b - \gamma b^T b = 0 \\ &\rightarrow \gamma = a^T b / b^T b \end{aligned}$$

**5.6 Gram-Schmidt algorithm.** Consider the list of  $n$   $n$ -vectors

$$a_1 = \begin{bmatrix} 1 \\ 0 \\ 0 \\ \vdots \\ 0 \end{bmatrix}, \quad a_2 = \begin{bmatrix} 1 \\ 1 \\ 0 \\ \vdots \\ 0 \end{bmatrix}, \quad \dots, \quad a_n = \begin{bmatrix} 1 \\ 1 \\ 1 \\ \vdots \\ 1 \end{bmatrix}.$$

(The vector  $a_i$  has its first  $i$  entries equal to one, and the remaining entries zero.) Describe what happens when you run the Gram-Schmidt algorithm on this list of vectors, i.e., say what  $q_1, \dots, q_n$  are. Is  $a_1, \dots, a_n$  a basis?

$$\begin{aligned} \tilde{f}_1 &= \vec{a}_1 \rightarrow f_1 = e_1 \\ \tilde{f}_2 &= \vec{a}_2 - (\vec{f}_1^T \vec{a}_2) \vec{f}_1 = \begin{bmatrix} 1 \\ 1 \\ 0 \\ \vdots \\ 0 \end{bmatrix} - \begin{bmatrix} 1 \\ 0 \\ 0 \\ \vdots \\ 0 \end{bmatrix} = \begin{bmatrix} 0 \\ 1 \\ 0 \\ \vdots \\ 0 \end{bmatrix} = e_2 \rightarrow f_2 = e_2 \\ \tilde{f}_3 &= \vec{a}_3 - (\vec{f}_1^T \vec{a}_3) \vec{f}_1 - (\vec{f}_2^T \vec{a}_3) \vec{f}_2 = \begin{bmatrix} 1 \\ 1 \\ 1 \\ \vdots \\ 0 \end{bmatrix} - \begin{bmatrix} 1 \\ 0 \\ 0 \\ \vdots \\ 0 \end{bmatrix} - \begin{bmatrix} 0 \\ 1 \\ 0 \\ \vdots \\ 0 \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ 1 \\ \vdots \\ 0 \end{bmatrix} = e_3 \\ &\rightarrow f_i = e_i \quad \text{basis } ?! \rightarrow \beta_1 a_1 + \dots + \beta_k a_k = 0 \rightarrow \text{because } a_n = \begin{bmatrix} 1 \\ 1 \\ \vdots \\ 1 \end{bmatrix} \text{ then} \\ \beta_n &= 0 \rightarrow \text{and } a_{n-1} = \begin{bmatrix} 1 \\ 1 \\ \vdots \\ 0 \end{bmatrix} \rightarrow \beta_{n-1} = 0 \rightarrow \text{so } a \text{ is basis} \end{aligned}$$

**5.7 Running Gram-Schmidt algorithm twice.** We run the Gram-Schmidt algorithm once on a given set of vectors  $a_1, \dots, a_k$  (we assume this is successful), which gives the vectors  $q_1, \dots, q_k$ . Then we run the Gram-Schmidt algorithm on the vectors  $q_1, \dots, q_k$ , which produces the vectors  $z_1, \dots, z_k$ . What can you say about  $z_1, \dots, z_k$ ?

since if vector is already orthogonal then applying Gram-Schmidt does nothing and it gives us:  $z_i = f_i$

**5.8** Early termination of Gram–Schmidt algorithm. When the Gram–Schmidt algorithm is run on a particular list of 10 15-vectors, it terminates in iteration 5 (since  $\tilde{q}_5 = 0$ ). Which of the following must be true?

- (a)  $a_2, a_3, a_4$  are linearly independent.
- (b)  $a_1, a_2, a_5$  are linearly dependent.
- (c)  $a_1, a_2, a_3, a_4, a_5$  are linearly dependent. ✓
- (d)  $a_4$  is nonzero.

**5.9** A particular computer can carry out the Gram–Schmidt algorithm on a list of  $k = 1000$   $n$ -vectors, with  $n = 10000$ , in around 2 seconds. About how long would you expect it to take to carry out the Gram–Schmidt algorithm with  $\tilde{k} = 500$   $\tilde{n}$ -vectors, with  $\tilde{n} = 1000$ ?

$$k = 1000, n = 10000 \rightarrow A_{10000 \times 1000} \rightarrow 2 \text{ sec}$$

book solution: comp speed  $\rightarrow$  flop/sec

Gram–schmidt requires  $2nk^2$  flops

$$k = 10^3, n = 10^4 \rightarrow \approx \frac{2 \times 10^{10} \text{ flop}}{2 \text{ sec}} = 10^{10} \text{ flop/sec}$$

second prob:  $k = 500, n = 10^3 \rightarrow 2 \times 10^3 \times 250000 \text{ flops} = 5 \times 10^8 \text{ flop}$

$$\rightarrow \frac{5 \times 10^8}{10^{10}} = 0.05 \text{ second}$$

**Calculate the weighted linear combination using Python code.**

$$\lambda_1 = 1, \lambda_2 = 2, \lambda_3 = -3$$

$$v_1 = \begin{bmatrix} 4 \\ 5 \\ 1 \end{bmatrix}, v_2 = \begin{bmatrix} -4 \\ 0 \\ -4 \end{bmatrix}, v_3 = \begin{bmatrix} 1 \\ 3 \\ 2 \end{bmatrix}$$

$$w = \lambda_1 v_1 + \lambda_2 v_2 + \lambda_3 v_3 = \begin{bmatrix} -7 \\ -4 \\ -13 \end{bmatrix}$$

```
import numpy as np
v1 = np.array([4, 5, 1])
v2 = np.array([-4, 0, -4])
v3 = np.array([1, 3, 2])
lambda1, lambda2, lambda3 = 1, 2, -3
weighted_sum = lambda1 * v1 + lambda2 * v2 + lambda3 * v3
print(f"weighted sum: {weighted_sum}")
```

**Run the following code multiple times and describe your results and observations.**

```
nPerClust = 50
# blur around centroid (std units)
blur = 1
# XY centroid locations
A = [ 1, 1 ]
B = [ -3, 1 ]
C = [ 3, 3 ]
# generate data
a = [ A[0]+np.random.randn(nPerClust)*blur ,
A[1]+np.random.randn(nPerClust)*blur ]
b = [ B[0]+np.random.randn(nPerClust)*blur ,
B[1]+np.random.randn(nPerClust)*blur ]
c = [ C[0]+np.random.randn(nPerClust)*blur ,
C[1]+np.random.randn(nPerClust)*blur ]
# concatinate into a matrix
data = np.transpose( np.concatenate((a,b,c),axis=1) )
# plot data
plt.plot(data[:,0],data[:,1],'ko',markerfacecolor='w')
plt.title('Raw (preclustered) data')
plt.xticks([])
plt.yticks([])
plt.show()

## initialize random cluster centroids
k = 3 # extract three clusters
```

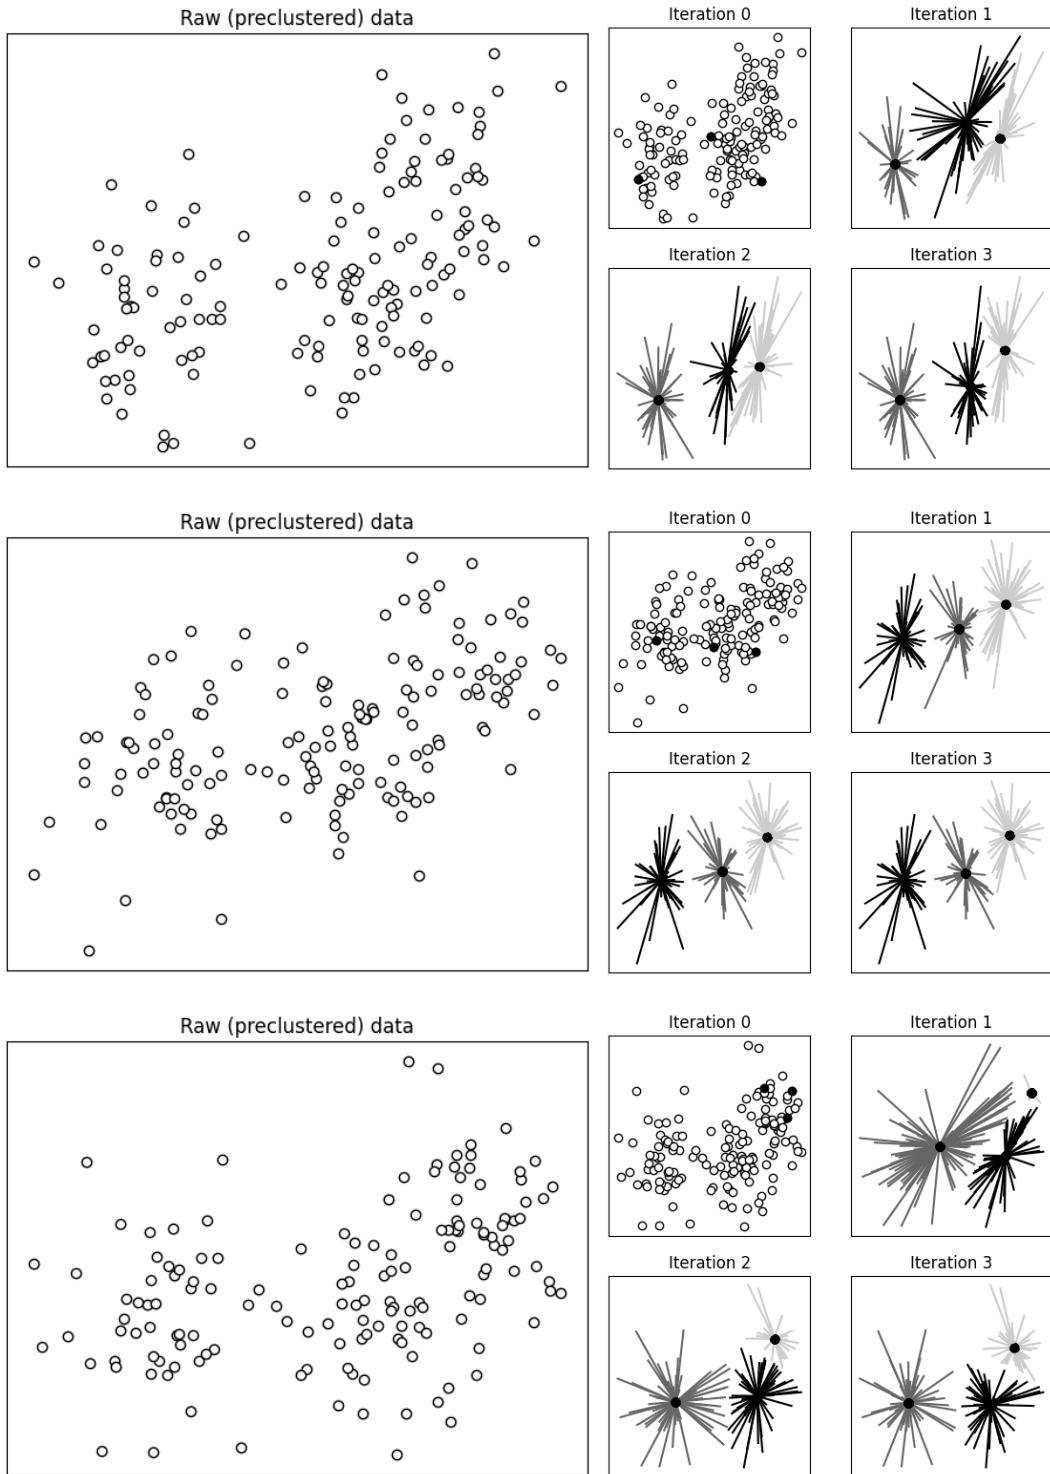
```

# random cluster centers (randomly sampled data points)
ridx = np.random.choice(range(len(data)),k,replace=False)
centroids = data[ridx,:]
# setup the figure
fig,axs = plt.subplots(2,2,figsize=(6,6))
axs = axs.flatten()
lineColors = [ [0,0,0],[.4,.4,.4],[.8,.8,.8] ]#'rbm'
# plot data with initial random cluster centroids
axs[0].plot(data[:,0],data[:,1],'ko',markerfacecolor='w')
axs[0].plot(centroids[:,0],centroids[:,1],'ko')
axs[0].set_title('Iteration 0')
axs[0].set_xticks([])
axs[0].set_yticks([])
# loop over iterations
for iteri in range(3):
    # step 1: compute distances
    dists = np.zeros((data.shape[0],k))
    for ci in range(k):
        dists[:,ci] = np.sum((data-centroids[ci,:])**2, axis=1)
    # step 2: assign to group based on minimum distance
    groupidx = np.argmin(dists, axis=1)
    # step 3: recompute centers
    for ki in range(k):
        centroids[ki,:] = [ np.mean(data[groupidx==ki,0]),
        np.mean(data[groupidx==ki,1]) ]

    # plot data points
    for i in range(len(data)):
        axs[iteri+1].plot([ data[i,0],centroids[groupidx[i],0] ],
        [ data[i,1],centroids[groupidx[i],1] ],color=lineColors[groupidx[i]])
    axs[iteri+1].plot(centroids[:,0],centroids[:,1],'ko')
    axs[iteri+1].set_title(f'Iteration {iteri+1}')
    axs[iteri+1].set_xticks([])
    axs[iteri+1].set_yticks([])
plt.show()

```

The given code implements the k-means clustering algorithm and visualizes the clustering process for a 2D dataset over four iterations. Here's a breakdown of the observations and results when running the code multiple times:



## K-means Clustering on Handwritten Digits (1-5) Dataset

This project explores the application of the K-means clustering algorithm to group handwritten digits from 1 to 5 using the MNIST dataset. The digits were extracted, flattened, and normalized before running the clustering algorithm. The goal was to cluster the data for varying values of  $k = 3, 4, 5, 6$  and visualize the resulting clusters and centroids.

### Data Preprocessing

- **Dataset:** The MNIST dataset was loaded using torchvision, which provides images of handwritten digits. Only the digits 1 through 5 were selected from the dataset.
- **Normalization:** Each image was flattened into a 28x28 (784) vector and normalized by dividing pixel values by 255.
- **GPU Acceleration:** The clustering algorithm leverages PyTorch to compute Euclidean distances and perform clustering operations on the GPU for faster processing.

### K-means Implementation

The K-means algorithm was implemented using PyTorch. The steps involved include:

1. **Initialization:** Centroids were randomly initialized from the dataset.
2. **Assignment:** Each data point was assigned to the nearest centroid based on the Euclidean distance.
3. **Update:** The centroids were updated by taking the mean of the points in each cluster.
4. **Convergence:** The algorithm iterates until the centroids no longer change or the maximum number of iterations is reached.

### Code Implementation

```
# Function to compute Euclidean distance using PyTorch
def euclidean_distance_gpu(a, b):
    a = a.unsqueeze(1) # Shape: (n_samples, 1, n_features)
    b = b.unsqueeze(0) # Shape: (1, n_clusters, n_features)
    return torch.sum((a - b) ** 2, dim=2)

# K-means algorithm using PyTorch
def run_kmeans_gpu(data, k, max_iter=500):
    # Initialize centroids randomly
    centroids = data[torch.randperm(data.shape[0])[:k]]
    for _ in tqdm(range(max_iter), desc="Running k-means"):
        # Compute distances and assign labels
        distances = euclidean_distance_gpu(data, centroids)
        labels = torch.argmin(distances, dim=1)
        # Compute new centroids
        new_centroids = torch.stack([
            data[labels == i].mean(dim=0) if (labels == i).sum() > 0 else centroids[i]
            for i in range(k)
        ])
```

```

# Check for convergence      if torch.allclose(new_centroids, centroids, atol=1e-8):
    print("K-means converged.")
    break
centroids = new_centroids
return centroids.cpu().numpy(), labels.cpu().numpy()

# Test K-means with multiple cluster sizes
k_to_test = [3, 4, 5, 6]
for k in k_to_test:
    print(f"Running k-means with k={k}...")
    centroids, labels = run_kmeans_gpu(x_normalized.cuda(), k)
    visualize_clusters(centroids, labels, x, k)
    print(f"Finished k-means with k={k}.")

```

## Visualization of Results

For each k value, the following visualizations were generated:

1. **Centroids:** Representative centroids for each cluster were displayed as images, reshaped into 28x28 matrices.
2. **Cluster Samples:** Random samples from each cluster were displayed in 28x28 image format to visually inspect the grouping.

## Results

For each k = 3, 4, 5, 6, the results were as follows:

- **Centroid Images:** The centroids for each cluster (shown as images) represented the "average" of the digits in that cluster. The centroids for different values of k displayed varying levels of differentiation, with more clusters yielding finer distinctions between different handwritten styles.
- **Cluster Samples:** For each value of k, at least 5 sample images from each cluster were displayed. This provided insight into how well the algorithm grouped similar-looking digits.

## Conclusion

The K-means algorithm was successful in clustering the digits 1 to 5 into meaningful groups. The quality of the clusters improved with higher values of k, although some overlap in cluster membership was observed when using fewer clusters. This demonstrates how K-means can be used for unsupervised learning tasks in digit classification.

## Example Visualizations

### Centroid and Sample Images:

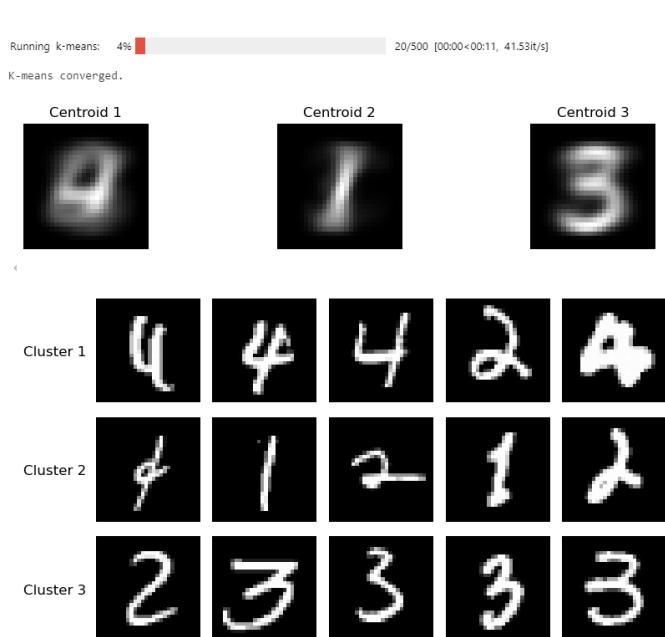


Figure 3-  $K = 3$

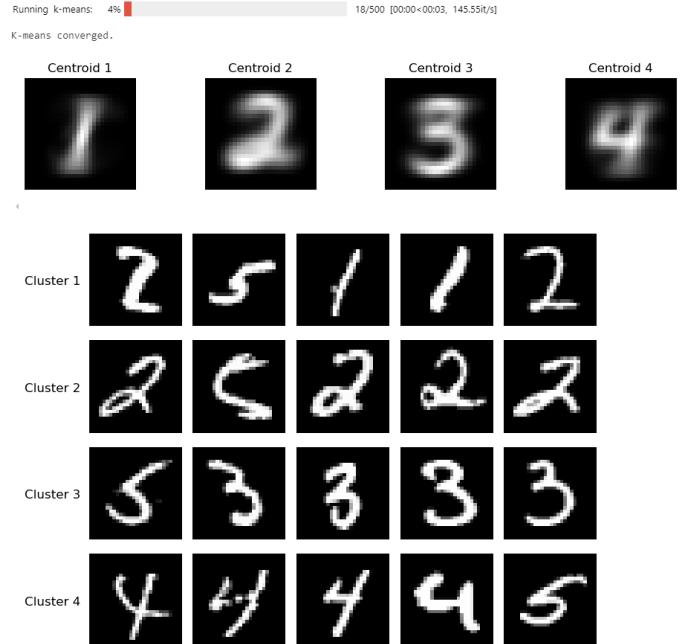


Figure 2-  $K = 4$

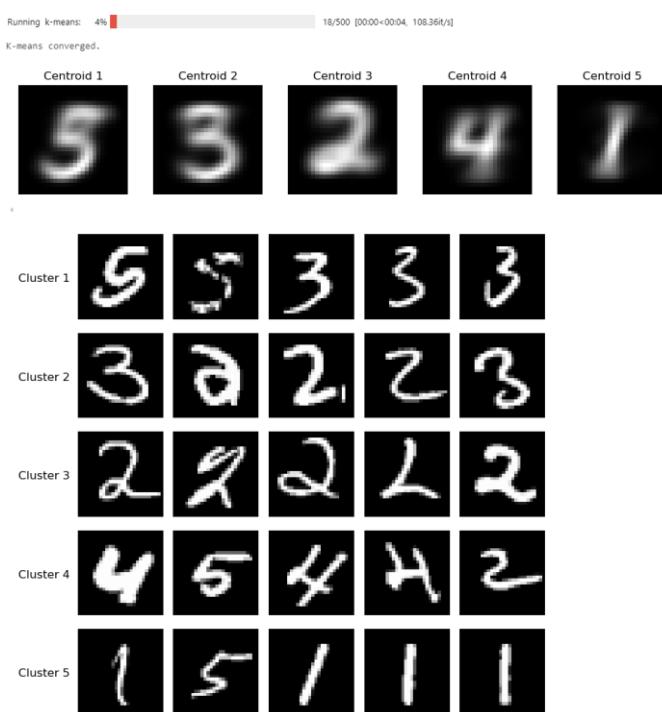


Figure 4-  $K = 5$

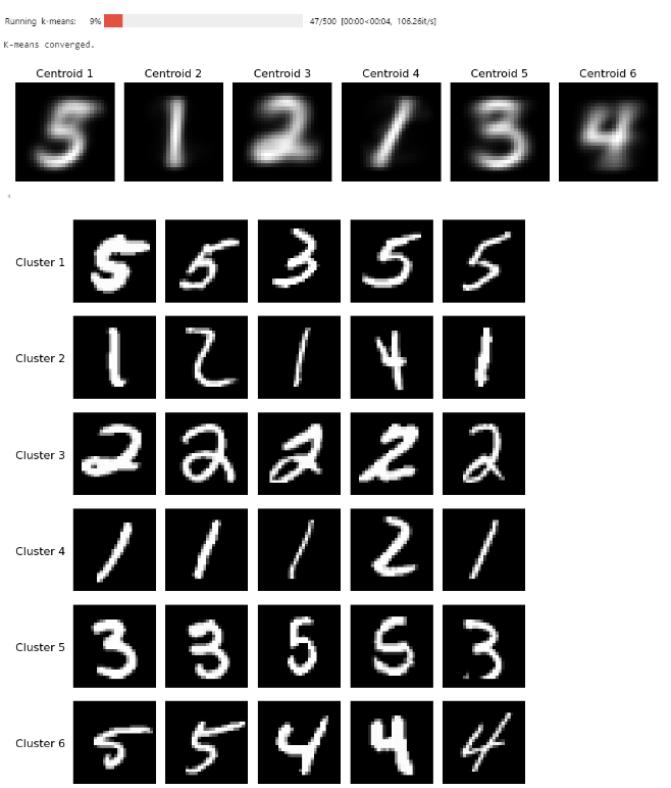


Figure 1-  $K = 6$